G R A P H P U L S E

Smarter Graphs, Faster Insights

A high-performance system that analyzes massive graph networks and learns from them — automatically.

C++ Core Engine | CPU Parallelism | GPU Acceleration | Machine Learning

Built for scale. Designed for results.

THE PROBLEM

The World Runs on Graphs

Massive Scale

Social networks, web links, and citation databases contain millions of interconnected nodes — impossible to process one-by-one.

Slow Analysis

Traditional sequential methods take hours on large graphs. In the real world, insights need to be fast and actionable.

Hidden Patterns

Valuable information is buried in the structure of connections — who knows who, which paper cites which — hard to surface without smart tools.

GraphPulse was built to solve exactly these challenges.

OUR SOLUTION

GraphPulse: One Pipeline, End-to-End

What GraphPulse Does

1. Takes in any large graph dataset
2. Automatically extracts meaningful features from the structure
3. Uses CPU and GPU together for maximum speed
4. Trains a machine learning model on the extracted features
5. Outputs predictions and performance metrics

Why It Matters

- ✓ Scales to graphs with millions of nodes
- ✓ Runs dramatically faster with hybrid parallelism
- ✓ Achieves over 86% accuracy on classification tasks
- ✓ Modular design — easy to extend or adapt
- ✓ Works on real-world datasets out of the box

HOW IT WORKS

The 9-Step Pipeline — Raw Data to Insight

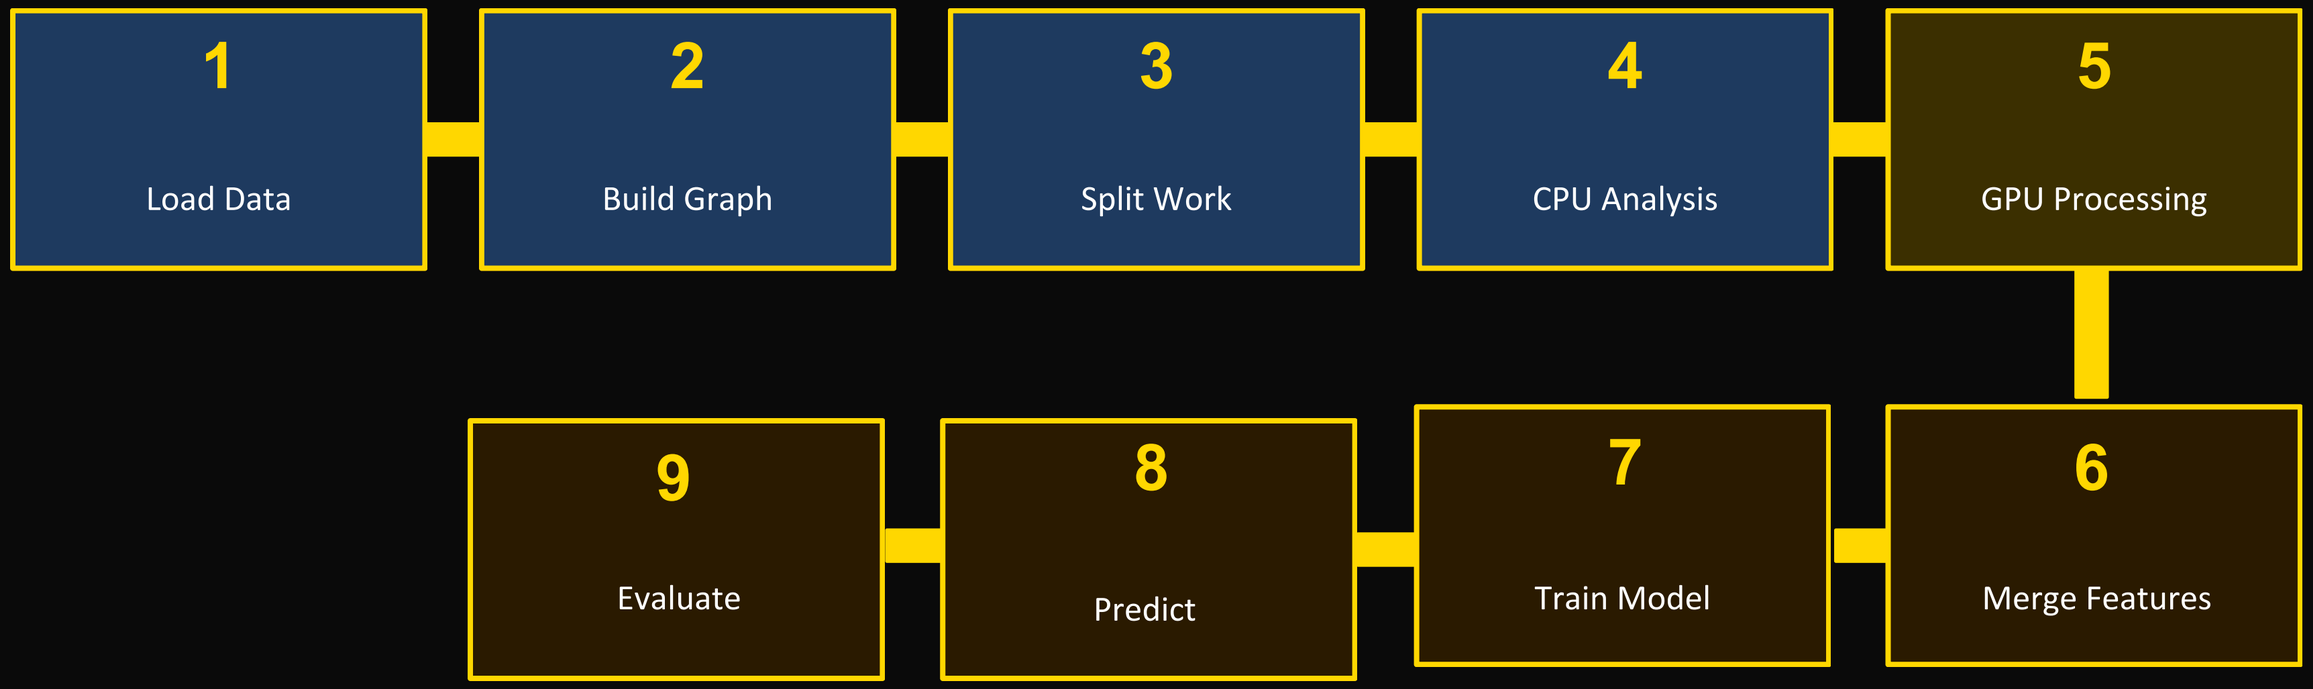

CPU lanes (steps 1–4) *GPU lanes (step 5)* *ML lanes (steps 7–9)*

Each step is independent and communicates through a standard data format — making the system easy to maintain and extend.

CONTEXT

What Is a Graph — and Why Analyze It?

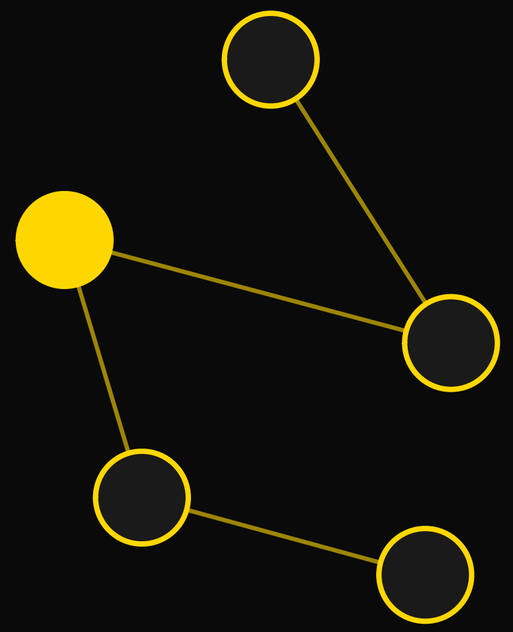

A network of nodes
& connections

Social Networks

People are nodes; friendships are edges. Analyzing structure reveals influencers, communities, and trends.

Web & Citation Graphs

Websites or papers linking to each other. Graph analysis powers search rankings and research discovery.

What We Extract

We compute features like 'how connected is a node?' or 'how central is it?' — then feed these to machine learning.

TECHNOLOGY

The Power of CPU + GPU Working Together

CPU — The Expert Analyst

Think of the CPU as a team of highly skilled experts. Each expert handles one complex task at a time — deeply and intelligently.

In GraphPulse, the CPU handles irregular, complex graph calculations like finding the most "connected" or "central" nodes (betweenness centrality).

We use a technology called OpenMP to split this work across multiple CPU cores simultaneously — like having many experts working in parallel.

VS

GPU — The Speed Powerhouse

Think of the GPU as thousands of simple workers doing the same job at once — extremely fast.

In GraphPulse, the GPU handles uniform, repetitive operations — like normalizing thousands of feature values at the same time.

We use CUDA (NVIDIA's GPU programming platform) to coordinate these thousands of parallel threads, giving us enormous throughput for batch operations.

GraphPulse intelligently routes each task to the right processor.

FEATURES

What We Measure in Every Graph

Degree Centrality

How many connections does a node have?

Like counting how many friends someone has on a social network. Nodes with many connections are "hubs".

Betweenness Centrality

How often does a node sit on the shortest path between others?

Like identifying the key bridge in a city road network — remove it and everything slows down.

PageRank

How important is a node based on who links to it?

The algorithm behind Google Search — a node is important if other important nodes point to it.

Clustering Coefficient

How tightly connected is a node's neighborhood?

Measures whether a node's neighbors also know each other — revealing tightly-knit communities.

MACHINE LEARNING

Teaching the System to Recognize Patterns

Think of it like a doctor diagnosing a patient: the 4 graph features are the symptoms, and the Random Forest model learns which combinations indicate which type of node.

01

Gather Features

For each node in the graph, we collect its 4 computed scores (degree, betweenness, pagerank, clustering).

02

Train the Forest

We run a Random Forest — hundreds of decision trees that vote together — on labeled training data to learn patterns.

86%+ Accuracy

03

Predict & Evaluate

On new, unseen data the model predicts node classes and we measure accuracy, F1 score, and use cross-validation.

RESULTS

What We Achieved

86%+	4	O(V·E)	3-Tier
ML Classification Accuracy	**Graph Features Extracted**	**Time Complexity**	**Parallel ML Strategy**
On real-world graph datasets	Per node, computed in parallel	Optimized with hybrid parallelism	Forest, tree, and node level

Tested On Real-World Datasets

SNAP CA-AstroPh · 4,039 nodes · 88,234 edges

SNAP CA-GrQc (Collaboration) · 5,242 nodes · 14,496 edges

Synthetic Barabási–Albert · 1K – 100K nodes · variable edges

ENGINEERING

Built for the Real World

Memory Safe

- No raw memory leaks — modern C++ RAII throughout
- GPU memory wrapped in safe resource handles
- Zero chance of silent crashes from dangling pointers

Highly Modular

- Each component connects through a clean interface
- Swap out any module without breaking the rest
- Interfaces: IGraph, IFeatureExtractor, IClassifier

Performance Tuned

- Adaptive scheduling — no wasted CPU cycles
- Pinned memory for fast CPU-GPU data transfers
- Lock-free parallel voting in Random Forest

Tested & Measured

- Benchmark scripts run across 1–8 thread configs
- Automated charts: speedup curves, accuracy vs size
- Deterministic seeding for reproducible results

TECHNOLOGY STACK

What Powers GraphPulse

Core Engine

C++ 17

The foundation. Fast, powerful systems language used for all pipeline logic and data structures.

CPU Parallelism

OpenMP

Automatically distributes graph analysis tasks across multiple CPU cores simultaneously.

GPU Acceleration

CUDA

NVIDIA's platform for running thousands of threads on a graphics card for massive throughput.

Machine Learning

Random Forest

An ensemble model that combines many decision trees to make accurate, robust predictions.

Build System

CMake

Manages the compilation process across different platforms and configurations.

Visualization

Python + Matplotlib

Scripts that automatically generate performance charts and accuracy graphs after each run.

CONCLUSION

GraphPulse — What We Built

1. A complete end-to-end system: from raw graph data to trained ML model and performance metrics — fully automated.
2. Hybrid CPU + GPU parallelism that scales to graphs with millions of nodes without sacrificing accuracy.
3. Over 86% classification accuracy using a parallelized Random Forest with corrected algorithms and real feature importance.
4. Production-quality code: modular interfaces, memory-safe design, benchmarking infrastructure, and automated visualization.

GraphPulse — Scalable · Intelligent · Fast